



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Chapter 7

MapReduce Types and Formats



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Learning Objectives

- MapReduce Types
- Input Formats
- Output Formats




Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

General Form of MapReduce

- General form:



- Combine function:

```
map: (K1, V1) → list(K2, V2)
combine: (K2, list(V2)) → list(K2, V2)
reduce: (K2, list(V2)) → list(K3, V3)
```



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Configuration of MapReduce Types – 1

Table 7-1. Configuration of MapReduce types in the new API

Property	Job setter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapreduce.job.inputformat.class	setInputFormatClass()	•	•				
mapreduce.map.output.key.class	setMapOutputKeyClass()			•			
mapreduce.map.output.value.class	setMapOutputValueClass()				•		
mapreduce.job.output.key.class	setOutputKeyClass()					•	
mapreduce.job.output.value.class	setOutputValueClass()						•
Properties that must be consistent with the types:							
mapreduce.job.map.class	setMapperClass()	•	•	•	•		
mapreduce.job.combine.class	setCombinerClass()			•	•		
mapreduce.job.partitioner.class	setPartitionerClass()			•	•		
mapreduce.job.output.key.comparator.class	setSortComparatorClass()			•			
mapreduce.job.output.group.comparator.class	setGroupingComparatorClass()			•			
mapreduce.job.reduce.class	setReducerClass()			•	•	•	•
mapreduce.job.outputformat.class	setOutputFormatClass()					•	•



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Configuration of MapReduce Types – 2

Table 7-2. Configuration of MapReduce types in the old API

Property	JobConf setter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapred.input.format.class	setInputFormat()	•	•				
mapred.mapoutput.key.class	setMapOutputKeyClass()			•			
mapred.mapoutput.value.class	setMapOutputValueClass()				•		
mapred.output.key.class	setOutputKeyClass()					•	
mapred.output.value.class	setOutputValueClass()						•
Properties that must be consistent with the types:							
mapred.mapper.class	setMapperClass()	•	•	•	•		
mapred.map.runner.class	setMapRunnerClass()	•	•	•	•		
mapred.combiner.class	setCombinerClass()			•	•		
mapred.partitionner.class	setPartitionerClass()			•	•		
mapred.output.key.comparator.class	setOutputKeyComparatorClass()			•			
mapred.output.value.groupfn.class	setOutputValueGroupingComparator()			•			
mapred.reducer.class	setReducerClass()			•	•	•	•
mapred.output.format.class	setOutputFormat()					•	•



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Without Setting a Mapper or a Reducer – 1

What happens when you run MapReduce without setting a mapper or a reducer? Let's try it by running this minimal MapReduce program:

```
public class MinimalMapReduce extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.printf("Usage: %s [generic options] <input> <output>\n",
                getClass().getSimpleName());
            ToolRunner.printGenericCommandUsage(System.err);
            return -1;
        }

        Job job = new Job(getConf());
        job.setJarByClass(getClass());
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduce(), args);
        System.exit(exitCode);
    }
}
```



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Without Setting a Mapper or a Reducer – 2

The only configuration that we set is an input path and an output path. We run it over a subset of our weather data with the following:

```
% hadoop MinimalMapReduce "input/ncdc/all/190{1,2}.gz" output
```

We do get some output: one file named *part-r-00000* in the output directory. Here's what the first few lines look like (truncated to fit the page):

```
0→00290290709999991901010106004+64333+023450FM-12+000599999V0202701N01591...  
0→00350290709999991902010106004+64333+023450FM-12+000599999V0201401N01181...  
135→00290290709999991901010113004+64333+023450FM-12+000599999V0202901N00821...  
141→00350290709999991902010113004+64333+023450FM-12+000599999V0201401N01181...  
270→00290290709999991901010120004+64333+023450FM-12+000599999V0209991C00001...  
282→00350290709999991902010120004+64333+023450FM-12+000599999V0201401N01391...
```

Each line is an integer followed by a tab character, followed by the original weather data record. Admittedly, it's not a very useful program, but understanding how it produces its output does provide some insight into the defaults that Hadoop uses when running MapReduce jobs. [Example 7-1](#) shows a program that has exactly the same



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

MapReduce Default Settings

Example 7-1. A minimal MapReduce driver, with the defaults explicitly set

```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {
```

```
    @Override
```

```
    public int run(String[] args) throws Exception {
        Job job = JobBuilder.parseInputAndOutput(this, getConf(), args);
        if (job == null) {
            return -1;
        }

```

```
        job.setInputFormatClass(TextInputFormat.class);
```

```
        job.setMapperClass(Mapper.class);
```

```
        job.setMapOutputKeyClass(LongWritable.class);
```

```
        job.setMapOutputValueClass(Text.class);
```

```
        job.setPartitionerClass(HashPartitioner.class);
```

```
        job.setNumReduceTasks(1);
```

```
        job.setReducerClass(Reducer.class);
```

```
        job.setOutputKeyClass(LongWritable.class);
```

```
        job.setOutputValueClass(Text.class);
```

```
        job.setOutputFormatClass(TextOutputFormat.class);
```

```
        return job.waitForCompletion(true) ? 0 : 1;
    }

```

```
    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }

```

```
}
```




Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Choosing the Number of Reducers

- The single reducer default is no good and needs to be set to a larger number.
- The optimal number of reducers is related to the total number of available reducer slots, which can be calculated as:
$$\text{Total nodes} * \text{tasktracker.reduce.tasks.maximum (default 2)}$$
- To have slightly fewer reducers than total slots.
 - Tolerate a few failures without extending job execution time.



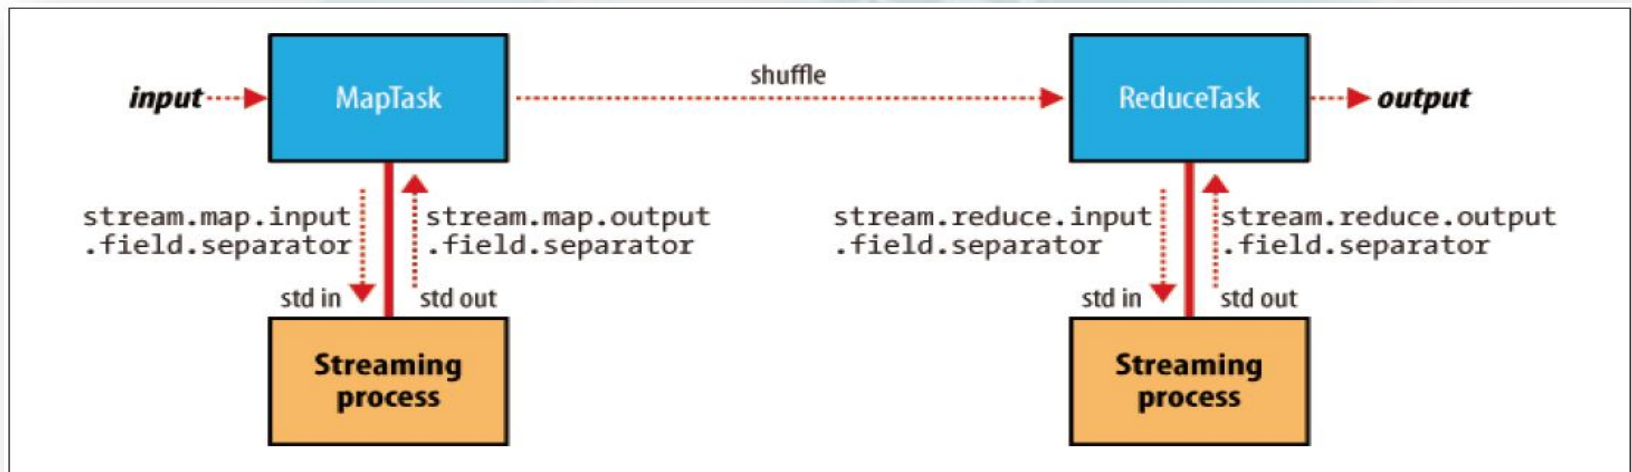
Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Keys and values in Streaming

A streaming application can control the separator.

- Default separator: tab character;
- Separators may be configured independently for maps and reduces;
- A key can be made up of the first n fields:
 - Ex) Output was a,b,c (and separator is a comma), n=2
 - Key: a,b Value:c





Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Input Splits

- **InputSplit**
 - A chunk of the input processed by a single map;
 - Each split is divided into records;
 - The map processes each record—a key-value pair;
 - Process the largest split first to minimize the job runtime;
 - The client sends the calculated splits to the JobTracker.



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

InputFormat Class Hierarchy

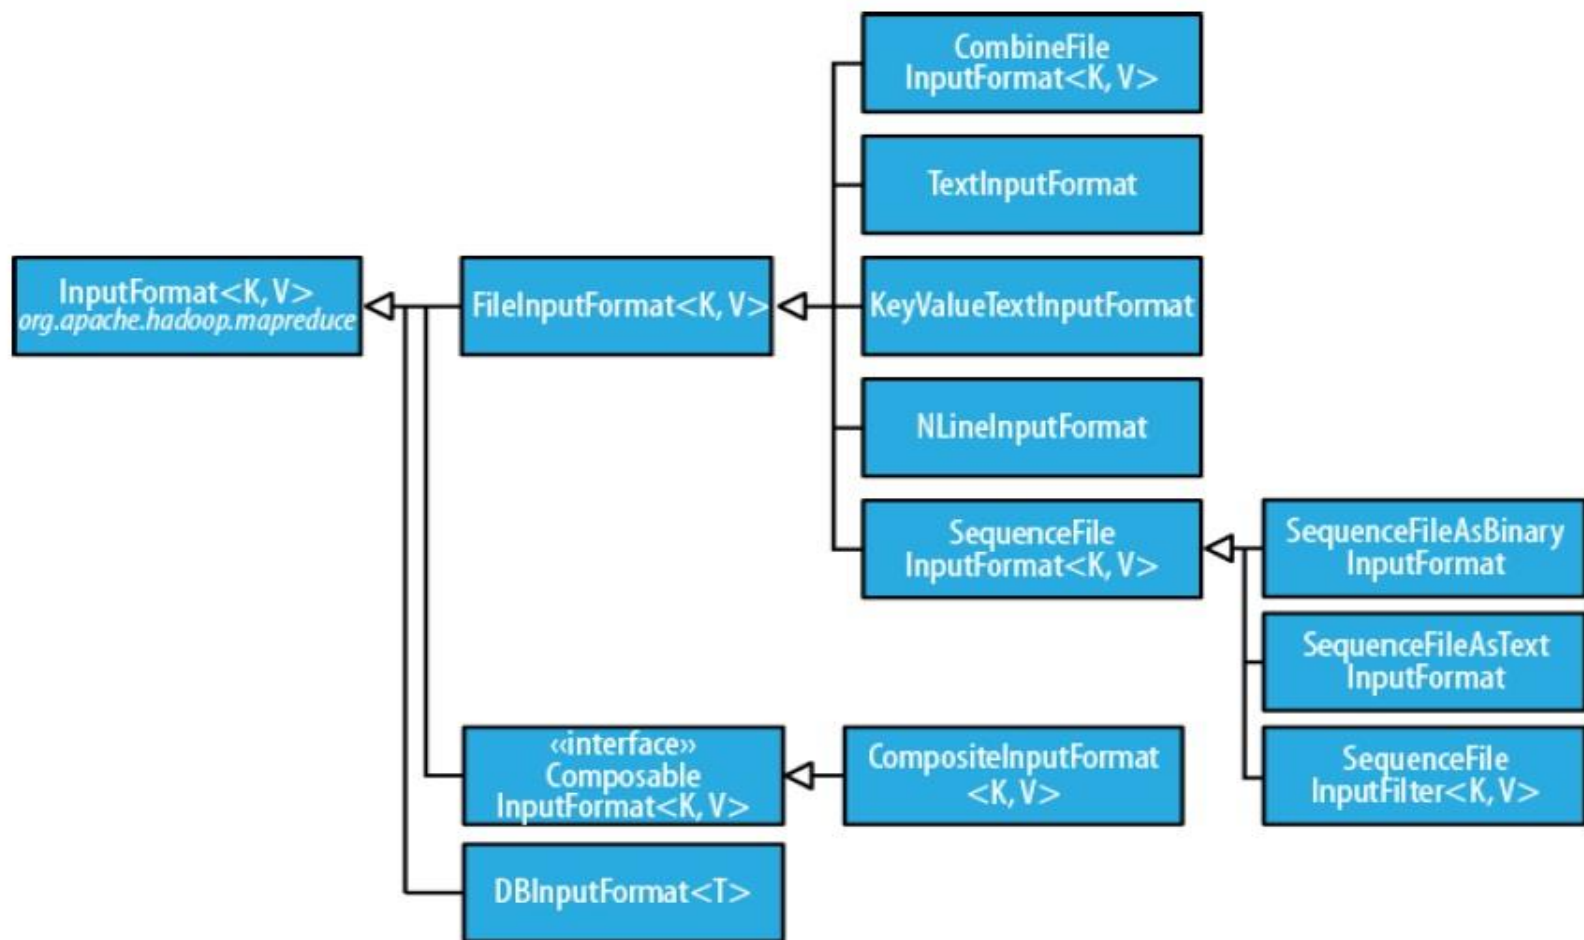


Figure 7-2. InputFormat class hierarchy



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Input Paths

FileInputFormat offers 4 static methods for setting a Job's input paths:

- `addInputPath()` and `addInputPaths()`
 - Add a path or paths to the list of inputs
 - Can call these methods repeatedly
- Two `setInputPaths()`
 - Set entire list of paths in one time (Replacing any paths that were set in previous calls)
- A path may represent
 - A file or a directory (consider all files in the directory as input)
 - A collection of files and directories by using a glob

```
public static void addInputPath(JobConf conf, Path path)
public static void addInputPaths(JobConf conf, String commaSeparatedPaths)
public static void setInputPaths(JobConf conf, Path... inputPaths)
public static void setInputPaths(JobConf conf, String commaSeparatedPaths)
```



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

FileInputFormat

- FileInputFormat uses a default filter that excludes hidden files.
- FileInputFormat splits only large files that larger than an HDFS block.
- The split size is normally the size of an HDFS block.
- The minimum split size is usually 1 byte.
- The maximum split size defaults to the maximum value of Java long type.



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Small Files and CombineFileInputFormat

- Hadoop works better with a small number of large files than a large number of small files.
 - FileInputFormat generates splits that each split is all or part of a single file
- Use CombineFileInputFormat to pack many files into splits.
 - Designed to work well with small files
 - Take node and rack locality when packing blocks into split
 - Worth when already have a large number of small files in HDFS
- Avoiding the many small files is a good idea.
 - Reduce the number of seeks
 - Merge small files into larger files by using a SequenceFile



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Preventing Splitting

- Some application don't want files to be split.
 - Want to process entire data by a single mapper
- Two ways of ensuring an existing file is not split:
 - Set the minimum split size larger than the largest file size
 - Override the `isSplittable()` method to return **false**



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Processing a Whole File As a Record – p.242

Example 7-2. An InputFormat for reading a whole file as a record

```
public class WholeFileInputFormat
    extends FileInputFormat<NullWritable, BytesWritable> {

    @Override
    protected boolean isSplittable(JobContext context, Path file) {
        return false;
    }

    @Override
    public RecordReader<NullWritable, BytesWritable> createRecordReader(
        InputSplit split, TaskAttemptContext context) throws IOException,
        InterruptedException {
        WholeFileRecordReader reader = new WholeFileRecordReader();
        reader.initialize(split, context);
        return reader;
    }
}
```



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

TextInputFormat – 1

- **TextInputFormat** is the default **InputFormat**
 - **Key**: The byte offset of the beginning of the line (**LongWritable**) ; Not line number
 - **Value**: The contents of the line excluding any line terminators (**Text**)

```
On the top of the Crumpetty Tree  
The Quangle Wangle sat,  
But his face you could not see,  
On account of his Beaver Hat.
```



```
(0, On the top of the Crumpetty Tree)  
(33, The Quangle Wangle sat,)  
(57, But his face you could not see,)  
(89, On account of his Beaver Hat.)
```




Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

TextInputFormat – 2

- The offset within the file of each line is known by each split independently of the other splits;
- Each split knows the size of the preceding splits;
- Add previous split size onto the offsets within the split to produce a global file offset.





Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

NLineInputFormat

- **TextInputFormat, KeyValueTextInputFormat** – each mapper receives a variable number of lines of input.
- **NLineInputFormat** – receive a fixed number of lines of input.
 - N: The number of lines of input.
 - Control N in `Mapred.line.input.format.linespermap` property
 - With N set to one, each mapper receives exactly one line of input.



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Binary Input

- **SequenceFileInputFormat**
 - Hadoop's sequence file format stores sequences of binary key-value pairs
 - Data is splittable (Data has sync points)
 - Use SequenceFileInputFormat
- **SequenceFileAsTextInputFormat**
 - Convert the sequence file's keys and values to Text objects
 - Use toString() method
- **SequenceFileAsBinaryInputFormat**
 - Retrieve the sequence file's keys and values as opaque binary objects



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Multiple Inputs

- Use `MultipleInput` when have data sources that provide the same type of data but in different formats, for example:
 - One might be tab-separated plain text, the other a binary sequence file;
 - Need to be parsed differently;
 - Use different mappers;
 - The map outputs have the same types;
 - Reducers are not aware of the different mappers.

```
MultipleInputs.addInputPath(job, ncddcInputPath,  
    TextInputFormat.class, MaxTemperatureMapper.class);  
MultipleInputs.addInputPath(job, metOfficeInputPath,  
    TextInputFormat.class, MetOfficeMaxTemperatureMapper.class);
```



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Output Formats

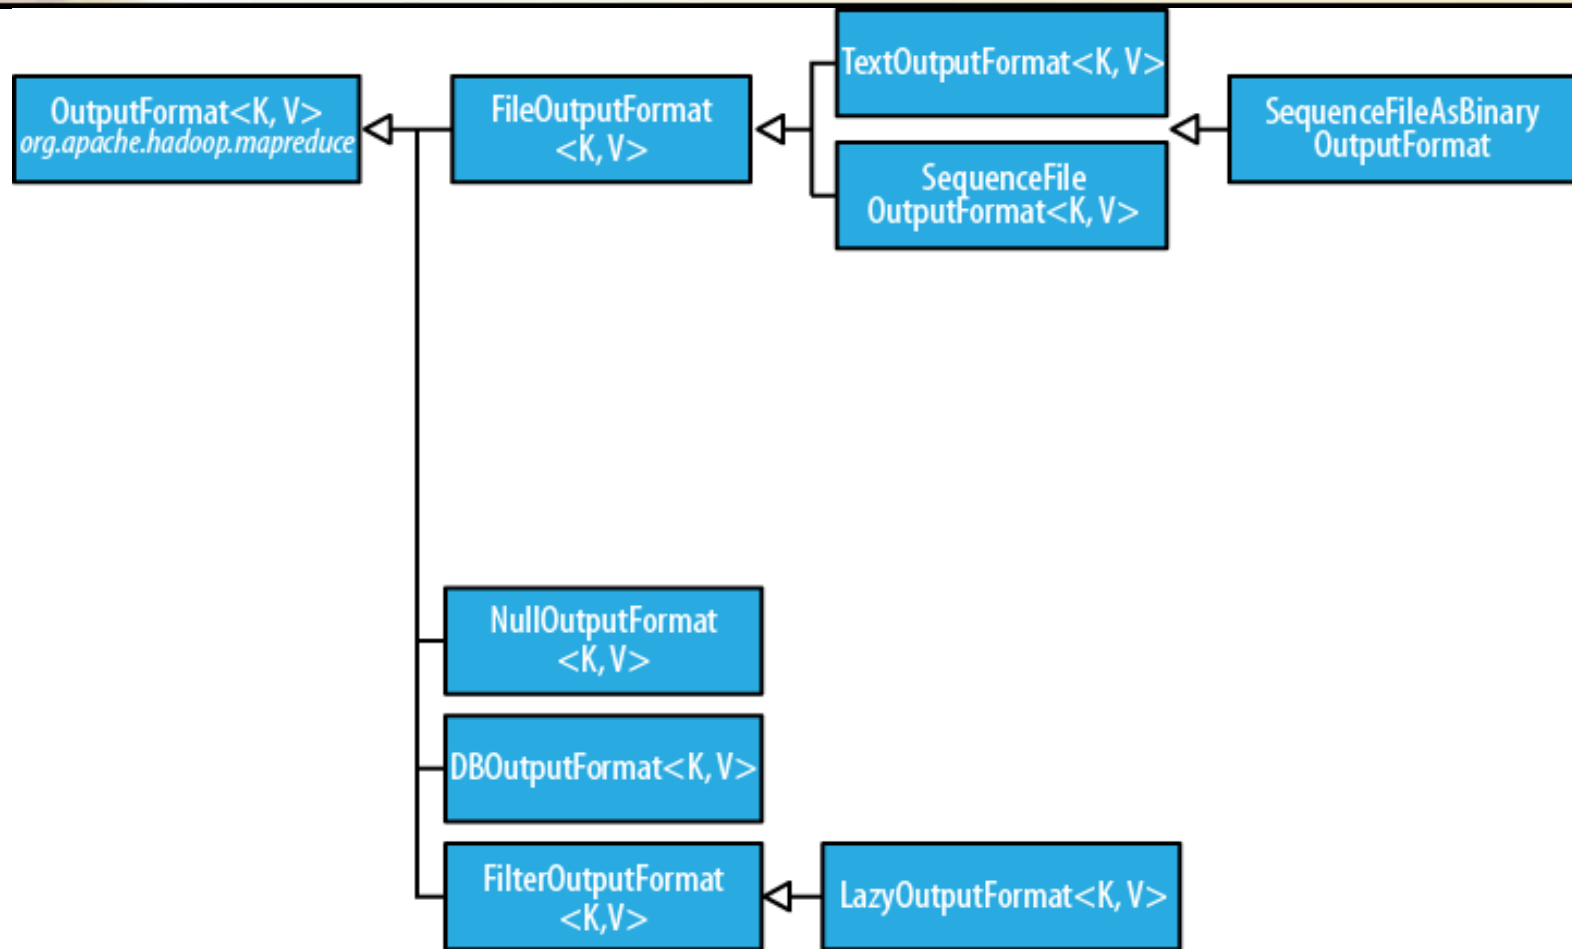


Figure 7-4. The OutputFormat class hierarchy



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Text Output

TextOutputFormat (default)

- Write records as lines of text
- Keys and Values may be of any type
 - It calls toString() method
- Each key-value pair is separated by a tab character
 - Set the separator in **mapred.textoutputformat.separator** property



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

Binary Output

- **SequenceFileOutputFormat**
 - Write sequence files for its output.
 - Compact, readily compressed (Useful for a further MapReduce job).
- **SequenceFileAsBinaryOutputFormat**
 - Write keys and values in raw binary format into a SequenceFile container.
- **MapFileOutputFormat**
 - Write MapFiles as output.
 - The keys in MapFile must be added in order.



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

MultipleOutputs – 1

MultipleOutputs

- Output file names are derived from the output keys and values, or from an arbitrary string.
- Allows each reducer (or mapper in a map-only job) to create more than a single file.
- Filenames are of the form name-m-nnnnn for map outputs and name-r-nnnnn for reduce outputs.
- Name is arbitrary and set by the program, and nnnnn is an integer designating the part number, starting from zero.



Hadoop – The Definitive Guide

Chapter 7: MapReduce Types and Formats

MultipleOutputs – 2

- P. 255 Example 7-5