



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Chapter 8

MapReduce Features



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Learning Objectives

- Counters
- Sorting
- Joins
- Side Data Distribution



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Counters

- Counters are a useful channel for gathering statistics about the job:
 - for quality control;
 - for application-level statistics;
 - for problem diagnosis (# of invalid records);
 - Easier to use and to retrieve than log output.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Built-In Counters

- Built-in counters for every job
 - Report various metrics for Map-Reduce jobs.
 - Counters for the number of bytes and records processed.

Table 8-1. Built-in counter groups

Group	Name/Enum	Reference
MapReduce task counters	<code>org.apache.hadoop.mapred.Task\$Counter (1.x)</code> <code>org.apache.hadoop.mapreduce.TaskCounter (post-1.x)</code>	Table 8-2
Filesystem counters	<code>FileSystemCounters (1.x)</code> <code>org.apache.hadoop.mapreduce.FileSystemCounter (post 1.x)</code>	Table 8-3
FileInputFormat counters	<code>org.apache.hadoop.mapred.FileInputFormat\$Counter (1.x)</code> <code>org.apache.hadoop.mapreduce.lib.input.FileInputFormatCounter (post-1.x)</code>	Table 8-4
FileOutputFormat counters	<code>org.apache.hadoop.mapred.FileOutputFormat\$Counter (1.x)</code> <code>org.apache.hadoop.mapreduce.lib.output.FileOutputFormatCounter (post-1.x)</code>	Table 8-5
Job counters	<code>org.apache.hadoop.mapred.JobInProgress\$Counter (1.x)</code> <code>org.apache.hadoop.mapreduce.JobCounter (post-1.x)</code>	Table 8-6



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Task and Job Counters

- **Task counters** gather information about tasks over the course of their execution.
 - MAP_INPUT_RECORDS counter counts the input records read by each map task;
 - Then aggregates over all map tasks in a job.
- **Job counters** are maintained by the jobtracker (or application master in YARN).
 - Measure job-level statistics,
 - TOTAL_LAUNCHED_MAPS counts the number of map tasks that were launched (including tasks that failed).



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

User-Defined Java Counters – 1

- MapReduce allows user code to define a set of counters, which are then incremented as desired in the mapper or reducer.
- Counters are defined by a Java enum, which serves to group related counters.
 - The name of the enum is the group name.
 - The enum's fields are the counter names.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

User-Defined Java Counters – 2

- P. 264: Example 8-1



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Retrieving Counters – P. 267

Example 8-2. Application to calculate the proportion of records with missing temperature fields

```
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.*;

public class MissingTemperatureFields extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 1) {
            JobBuilder.printUsage(this, "<job ID>");
            return -1;
        }
        String jobID = args[0];
        JobClient jobClient = new JobClient(new JobConf(getConf()));
        RunningJob job = jobClient.getJob(JobID.forName(jobID));
        if (job == null) {
            System.err.printf("No job with ID %s found.\n", jobID);
            return -1;
        }
        if (!job.isComplete()) {
            System.err.printf("Job %s is not complete.\n", jobID);
            return -1;
        }

        Counters counters = job.getCounters();
        long missing = counters.getCounter(
            MaxTemperatureWithCounters.Temperature.MISSING);

        long total = counters.getCounter(Task.Counter.MAP_INPUT_RECORDS);

        System.out.printf("Records with missing temperature fields: %.2f%%\n",
            100.0 * missing / total);
        return 0;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MissingTemperatureFields(), args);
        System.exit(exitCode);
    }
}
```




Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Sorting

- The ability to sort data is at the heart of MapReduce.
- Even if your application isn't concerned with sorting per se, it may be able to use the sorting stage that MapReduce provides to organize its data.
- P. 269: Example 8-3.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Sorting a Sequence File with IntWritable – P. 270

Example 8-4. A MapReduce program for sorting a SequenceFile with IntWritable keys using the default HashPartitioner

```
public class SortByTemperatureUsingHashPartitioner extends Configured
    implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        Job job = JobBuilder.parseInputAndOutput(this, getConf(), args);
        if (job == null) {
            return -1;
        }

        job.setInputFormatClass(SequenceFileInputFormat.class);
        job.setOutputKeyClass(IntWritable.class);
        job.setOutputFormatClass(SequenceFileOutputFormat.class);
        SequenceFileOutputFormat.setCompressOutput(job, true);
        SequenceFileOutputFormat.setOutputCompressorClass(job, GzipCodec.class);
        SequenceFileOutputFormat.setOutputCompressionType(job,
            CompressionType.BLOCK);

        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new SortByTemperatureUsingHashPartitioner(),
            args);
        System.exit(exitCode);
    }
}
```



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Produce MapFiles as Output – 271

Example 8-5. A MapReduce program for sorting a SequenceFile and producing MapFiles as output

```
public class SortByTemperatureToMapFile extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        Job job = JobBuilder.parseInputAndOutput(this, getConf(), args);
        if (job == null) {
            return -1;
        }

        job.setInputFormatClass(SequenceFileInputFormat.class);
        job.setOutputKeyClass(IntWritable.class);
        job.setOutputFormatClass(MapFileOutputFormat.class);
        SequenceFileOutputFormat.setCompressOutput(job, true);
        SequenceFileOutputFormat.setOutputCompressorClass(job, GzipCodec.class);
        SequenceFileOutputFormat.setOutputCompressionType(job,
            CompressionType.BLOCK);

        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new SortByTemperatureToMapFile(), args);
        System.exit(exitCode);
    }
}
```




Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Retrieve the First Entry

Example 8-6. Retrieve the first entry with a given key from a collection of MapFiles

```
public class LookupRecordByTemperature extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            JobBuilder.printUsage(this, "<path> <key>");
            return -1;
        }
        Path path = new Path(args[0]);
        IntWritable key = new IntWritable(Integer.parseInt(args[1]));

        Reader[] readers = MapFileOutputFormat.getReaders(path, getConf());
        Partitioner<IntWritable, Text> partitioner =
            new HashPartitioner<IntWritable, Text>();
        Text val = new Text();
        Writable entry =
            MapFileOutputFormat.getEntry(readers, partitioner, key, val);
        if (entry == null) {
            System.err.println("Key not found: " + key);
            return -1;
        }
        NcdcRecordParser parser = new NcdcRecordParser();
        parser.parse(val.toString());
        System.out.printf("%s\t%s\n", parser.getStationId(), parser.getYear());
        return 0;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new LookupRecordByTemperature(), args);
        System.exit(exitCode);
    }
}
```



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Total Sort

- Produce a set of sorted files that, if concatenated, would form a globally sorted file.
 - Use a partitioner that respects the total order of the output.
- Although this approach works, you have to choose your partition sizes carefully to ensure that they are fairly even so that job times aren't dominated by a single reducer.

Temperature range	< -10°C	[-10°C, 0°C)	[0°C, 10°C)	>= 10°C
Proportion of records	11%	13%	17%	59%



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Sampling

- To construct more even partitions, we need to have a better understanding of the distribution for the whole dataset.
- It's possible to get a fairly even set of partitions, by *sampling* the key space.
- The idea behind sampling is that you look at a small subset of the keys to approximate key distribution, which is then used to construct partitions.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Global Sort

- P. 276: Example 8-8




Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Secondary Sort – 1

- Example: calculating the maximum temperature for each year.
- It is possible to impose an order on the values by sorting and grouping the keys in a particular way.
 - Make the key a composite of the natural key and the natural value.
 - The key comparator should order by the composite key, that is, the natural key and natural value.
 - The partitioner and grouping comparator for the composite key should consider only the natural key for partitioning and grouping.





Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Secondary Sort – 2

- P. 279: Example 8-9



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Map-Side Join – 1

- A map-side join works by performing the join before the data reaches the map function. Requirements:
 - Each input dataset must be divided into the same number of partitions.
 - It must be sorted by the same key (the join key) in each source.
- Above requirements actually fit the description of the output of a MapReduce job:
 - A map-side join can be used to join the outputs of several jobs that had the same number of reducers, the same keys, and output files that are not splittable.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Map-Side Join – 2

Stations

Station ID	Station Name
011990-99999	SIHCCAJAVRI
012650-99999	TYNSET-HANSMOEN

Records

Station ID	Timestamp	Temperature
012650-99999	194903241200	111
012650-99999	194903241800	78
011990-99999	195005150700	0
011990-99999	195005151200	22
011990-99999	195005151800	-11

Join

Station ID	Station Name	Timestamp	Temperature
011990-99999	SIHCCAJAVRI	195005150700	0
011990-99999	SIHCCAJAVRI	195005151200	22
011990-99999	SIHCCAJAVRI	195005151800	-11
012650-99999	TYNSET-HANSMOEN	194903241200	111
012650-99999	TYNSET-HANSMOEN	194903241800	78



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Reduce-Side Join – 1

- Reduce-side join is more general than a map-side join:
 - Input datasets don't have to be structured in any particular way.
 - Less efficient as both datasets have to go through the MapReduce shuffle.
 - The mapper tags each record with its source.
 - Uses the join key as the map output key so that the records with the same key are brought together in the reducer.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Reduce-Side Join – 2

Example 8-12. Mapper for tagging station records for a reduce-side join

```
public class JoinStationMapper
    extends Mapper<LongWritable, Text, TextPair, Text> {
    private NcdcStationMetadataParser parser = new NcdcStationMetadataParser();

    @Override
    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        if (parser.parse(value)) {
            context.write(new TextPair(parser.getStationId(), "0"),
                new Text(parser.getStationName()));
        }
    }
}
```

Example 8-13. Mapper for tagging weather records for a reduce-side join

```
public class JoinRecordMapper
    extends Mapper<LongWritable, Text, TextPair, Text> {
    private NcdcRecordParser parser = new NcdcRecordParser();

    @Override
    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        parser.parse(value);
        context.write(new TextPair(parser.getStationId(), "1"), value);
    }
}
```

Example 8-14. Reducer for joining tagged station records with tagged weather records

```
public class JoinReducer extends Reducer<TextPair, Text, Text, Text> {

    @Override
    protected void reduce(TextPair key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {
        Iterator<Text> iter = values.iterator();
        Text stationName = new Text(iter.next());
        while (iter.hasNext()) {
            Text record = iter.next();
            Text outValue = new Text(stationName.toString() + "\t" + record.toString());
            context.write(key.getFirst(), outValue);
        }
    }
}
```



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Side Data Distribution

- Side data are extra read-only data needed by a job to process the main dataset.
- The challenge is to make side data available to all the map or reduce tasks (which are spread across the cluster).
 - Using the Job Configuration
 - Distributed Cache



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Using the Job Configuration

- Set arbitrary key-value pairs in the job configuration using the various setter methods on **Configuration**.
- Useful if one needs to pass a small piece of metadata to tasks.
- Don't use this mechanism for transferring more than a few kilobytes of data.



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Using Distributed Cache – 1

- Distribute datasets using Hadoop's distributed cache mechanism.
- Provides a service for copying files and archives to the task nodes in time for the tasks to use them when they run.
- `GenericOptionsParser` – specify the files to be distributed as a comma-separated list of URIs as the argument to the **-files** option.

```
% hadoop jar job.jar MaxTemperatureByStationNameUsingDistributedCacheFile \  
-files input/ncdc/metadata/stations-fixed-width.txt input/ncdc/all output
```



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Using Distributed Cache – 2

- P. 290: Example 8-16



Hadoop – The Definitive Guide

Chapter 8: MapReduce Features

Using Distributed Cache – 3

- Hadoop copies the file specified by the `-file` and `-archives` options to the jobtracker's filesystem (normally HDFS).
- Before a task is run, the tasktracker copies the files from the jobtracker's filesystem to a local disk—the cache—so the task can access the files.
- The tasktracker also maintains a reference count for the number of tasks using each file in the cache.
- After the task has run, the file's reference count is decreased by one, and when it reaches zero it is eligible for deletion.
- Files are deleted to make room for a new file when the cache exceeds a certain size—10 GB by default.